



DC4Cities

Fabien Hermenier



Episode 3





the goal of DC4Cities is to let existing
and new data centres become energy
adaptive

adapt

manage the workload wrt.
renewable energy availability

be adapted

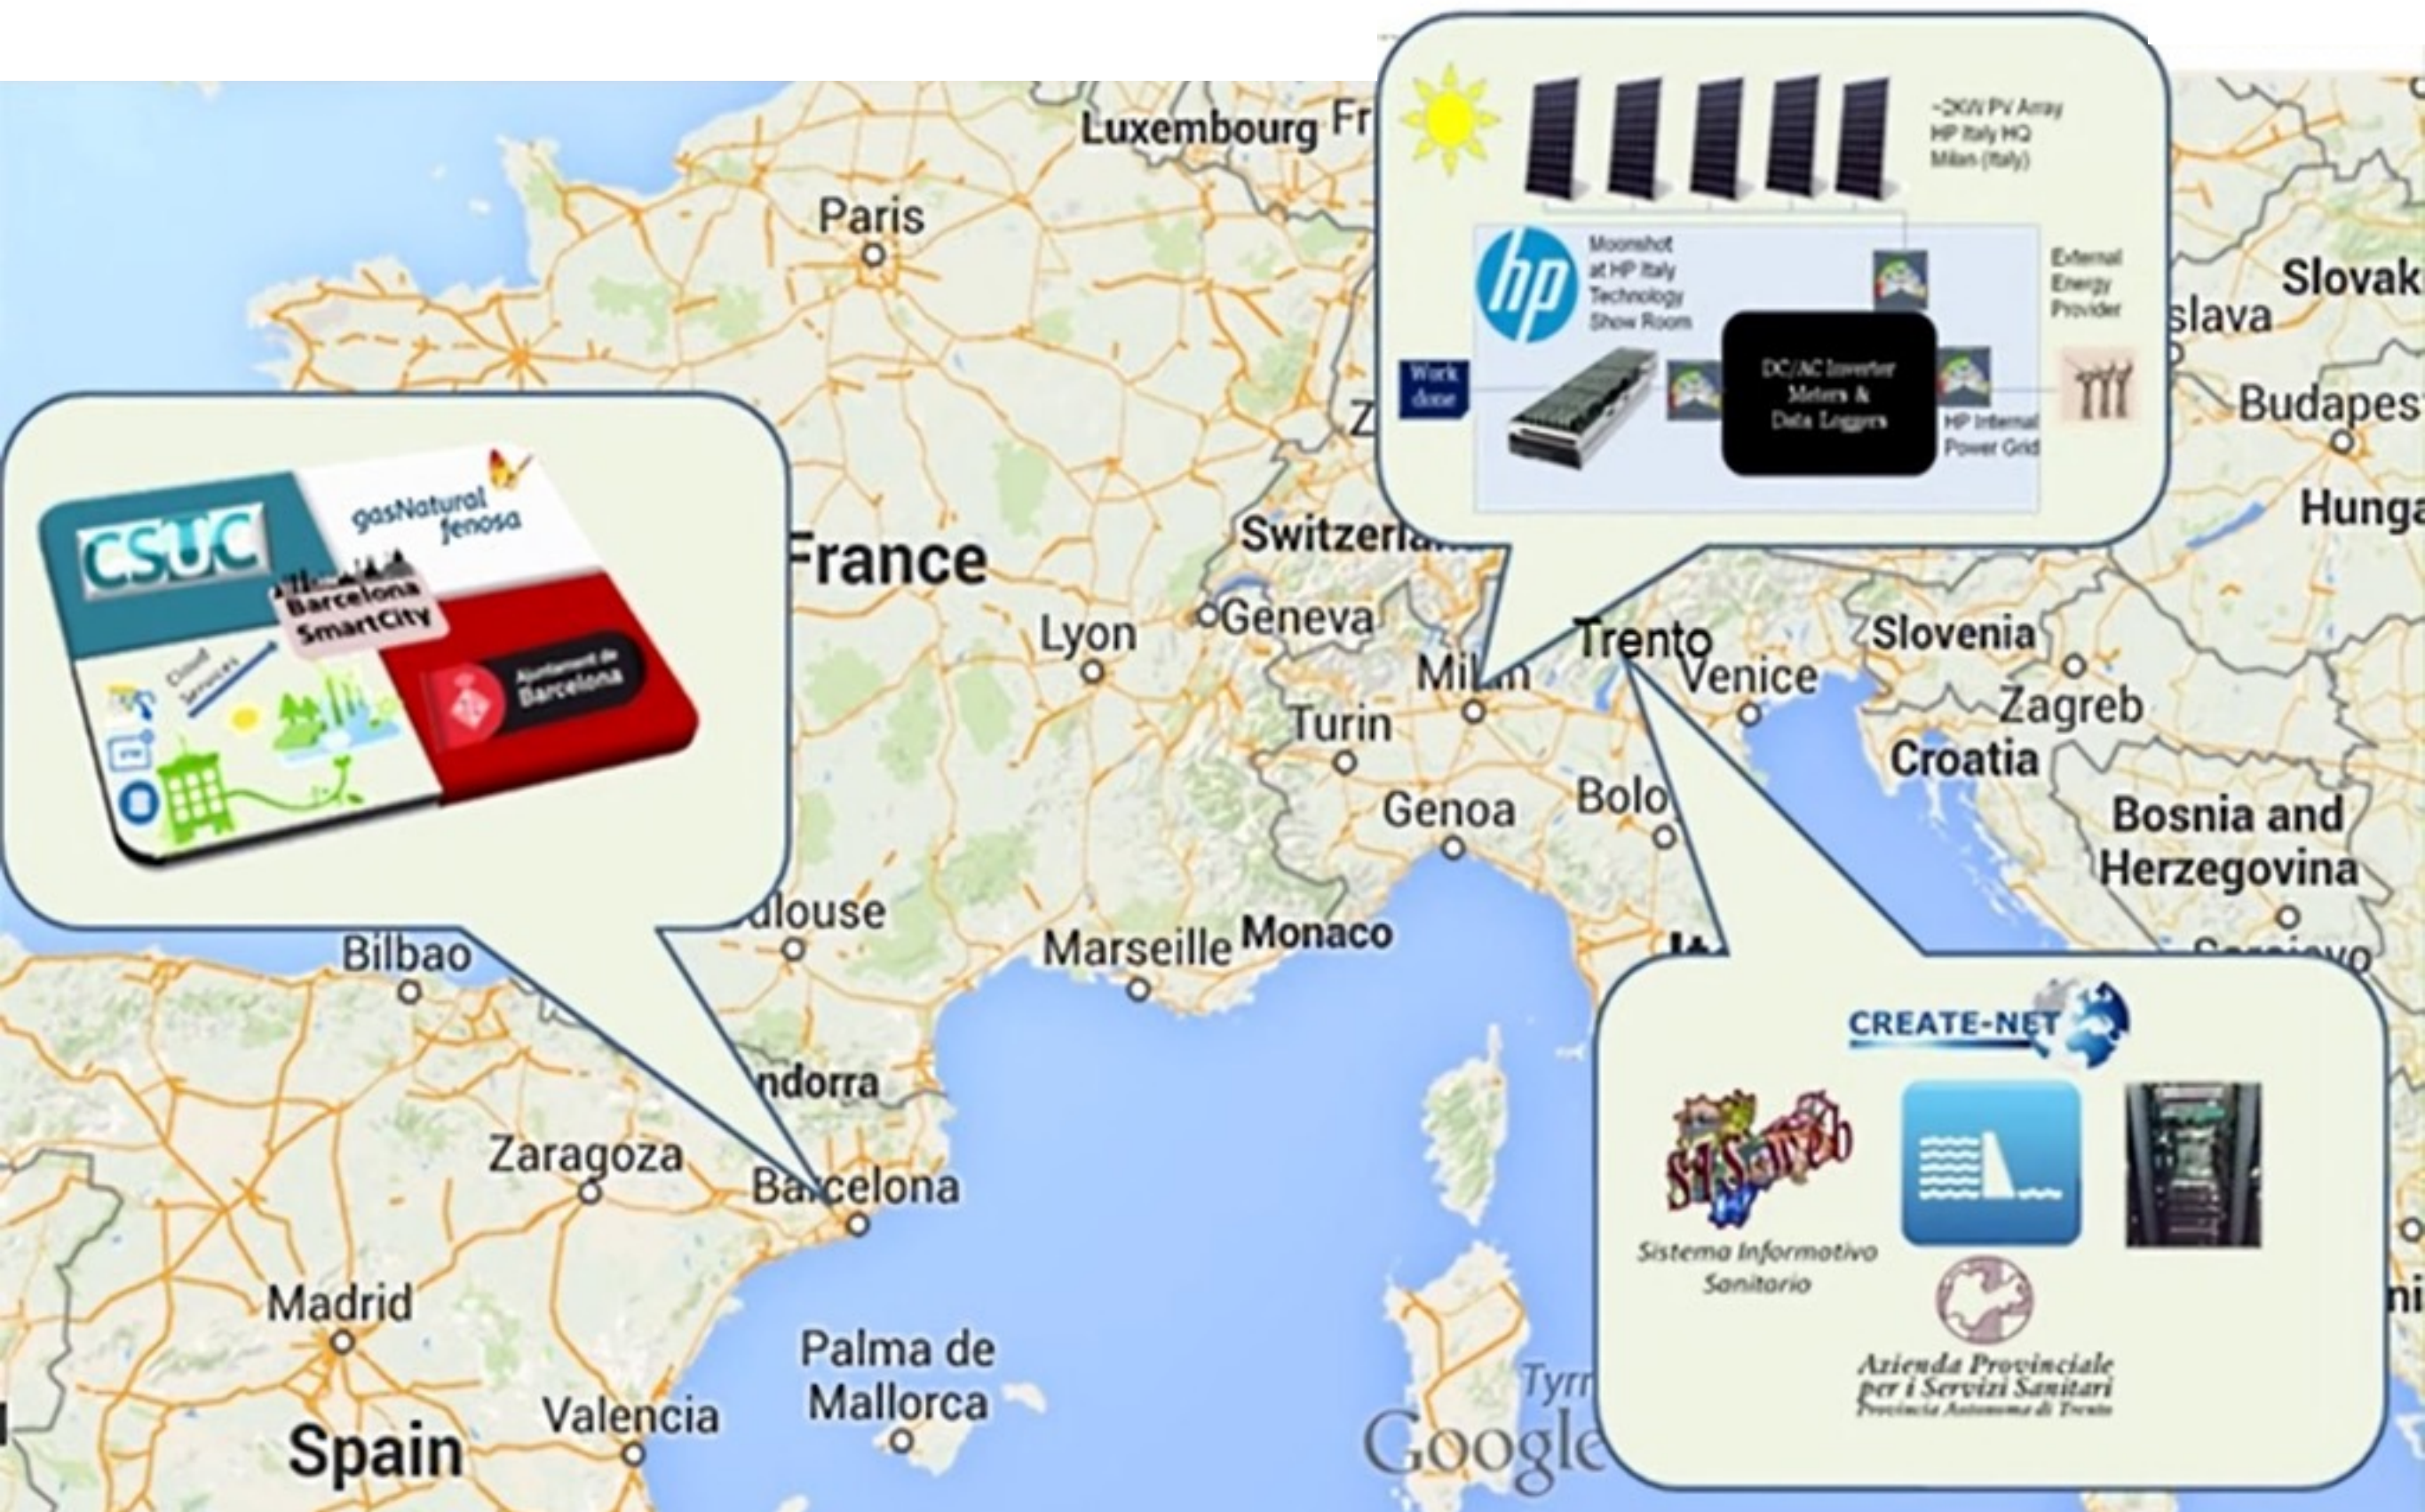
reply to a regulation of a smart
city energy authority



10 partners,
30 months



3 trial sites, 2 periods



TL;DR

metrics

energy modeling


adaptation mechanisms

Beyond PUE & CUE

abstract
not business oriented

Measuring the work done
Watt per *whatever* request

Measuring adaptivity
work shifting capabilities





manual modeling
expertise, validation, robustness issues

modeling power and performance using machine learning



- dominant variable detection
- model generation
- model validation
- output: functions, plots, ...

HP-trial PV production



$$\hat{y}(t) = -0.49 + \frac{1.49}{1 + e^{-0.0009 x_1(t) + 0.75}} + \frac{0.09}{1 + e^{-0.6 x_2(t) + 11.4}}$$

x1: irradiation
x2: outside temperature

CSUC VM power consumption



$$\hat{y}(t) = -343 + \frac{528}{1 + e^{-0.003 x_1(t) + 1.4}}$$

x1: vCPU usage

adapt

be adapted

Act 1

Energy Adaptive Software Components

attached to an application

exhibit

- working modes
- actuators

(see UCC'15 paper)

```
- name: pageIndexing  
  businessUnit: kPage
```


workingModes:

- ```
- name: WM0
 actuator: bin/run.sh WM0
 performance: !amount '0 kPage/h'
 power: !amount '6 W'
```
- 
- ```
- name: WM1  
  actuator: bin/run.sh WM1  
  performance: !amount '32 kPage/h'  
  power: !amount '27 W'
```
-
- ```
- name: WM2
 actuator: bin/run.sh WM2
 performance: !amount '60 kPage/h'
 power: !amount '33 W'
```




Each EASC proposes  
alternative execution plans  
matching a power budget


One consolidator chooses  
one plan for each EASC  
maximising ren%



greedy



proportional



aggressive

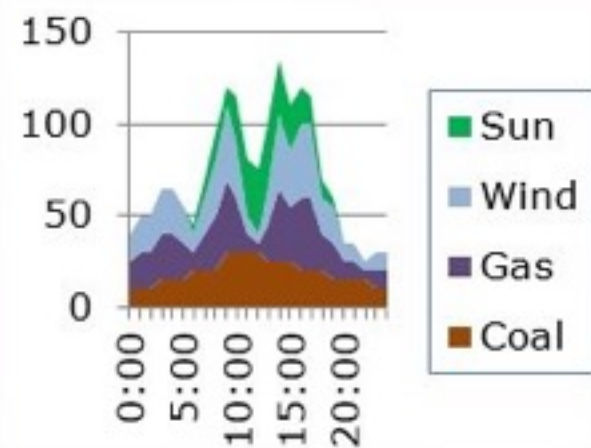
## Grid/Smart Grid



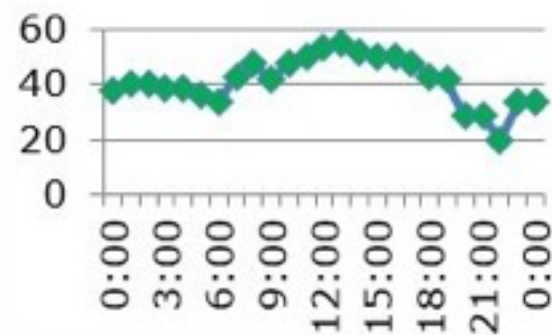
## Renewable Energy Providers



## Smart City Control



### RenPct




## Energy Availability Forecast

## Weather Forecast




### PV power



## Power/Energy Goals

**80%  
Ren**

## Renewable Energy Adaptive Interface





*Data Centre  
Energy  
Controller*




# Data Centre Energy Controller

## DC Ideal power




**Service  
Quota  
Split  
Policies**


## Quota B




## Quota A



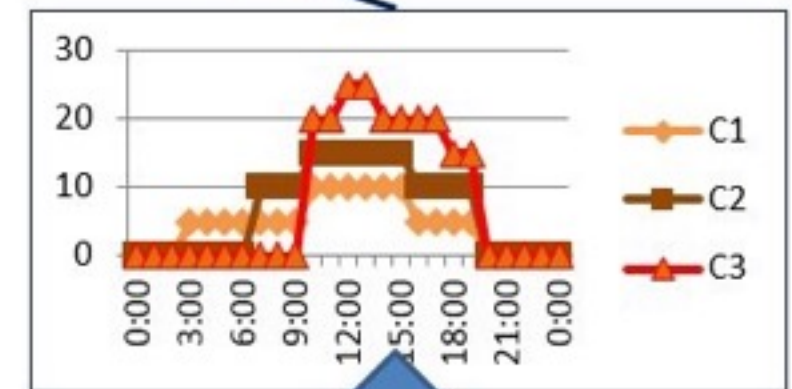
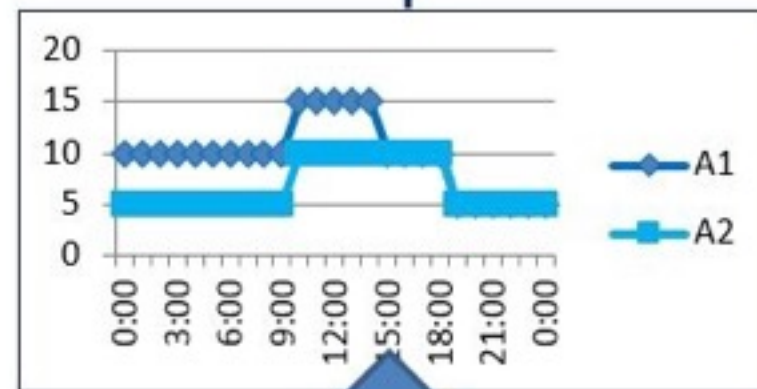


## Quota C




Energy Adaptive Data Centre Operation Interface



## Data Centre Energy Controller




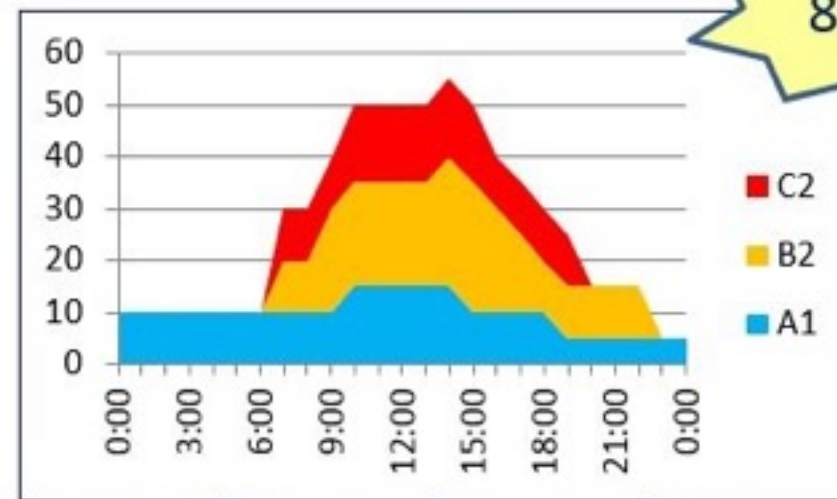


Energy Adaptive Data Centre Operation Interface






## Data Centre Energy Controller



Energy Adaptive Data Centre Operation Interface





# Lessons learned

## The good

it worked at M10  
quite scalable



## The bad

software && hardware dependent  
flexibility at the client side

## The ugly


some magic assumptions  
very limited flexibility



# Act 2

- The takeover -

# shape EASCs for sustainable profitability



pick WMs such as

$$\min(\text{penalty}(\text{SLO}) + \text{penalty}(\text{SMA}) + \text{price}(\text{E}))$$

# Energy Adaptive Software Components

attached to an application


exhibit

- working modes
- SLO (cumulative or instant)
- transition costs
- actuators

```
- name: pageIndexing
 businessUnit: kPage
 SLO:
 - timeFrom: 00:00:00
 timeTo: 24:00:00
 cumulativeObjective: !amount '200 kPage'
 basePrice: !amount '100 EUR'
 priceModifiers:
 - threshold: !amount '200 kPage'
 penalty: !amount '0 EUR/kPage'
 - threshold: !amount '100 kPage'
 penalty: !amount '-1 EUR/kPage'
 - threshold: !amount '0 kPage'
 penalty: !amount '-100 EUR'
 workingModes:
 - name: WM0
 actuator: bin/run.sh WM0
 performance: !amount '0 kPage/h'
 power: !amount '6 W'
 transitions:
 - target: WM1
 performanceCost: !amount '1 kPage'
 - target: WM2
 performanceCost: !amount '2 kPage'
 - name: WM1
 actuator: bin/run.sh WM1
 performance: !amount '32 kPage/h'
 power: !amount '27 W'
 transitions:
 - target: WM2
 performanceCost: !amount '2 kPage'
 - name: WM2
 actuator: bin/run.sh WM2
 performance: !amount '60 kPage/h'
 power: !amount '33 W'
```

EASC

weighted automata  
with counters





+ penalty functions for the Smart City Authority  
the SLA





# Hewlett Packard Enterprise

grid renewable  
part






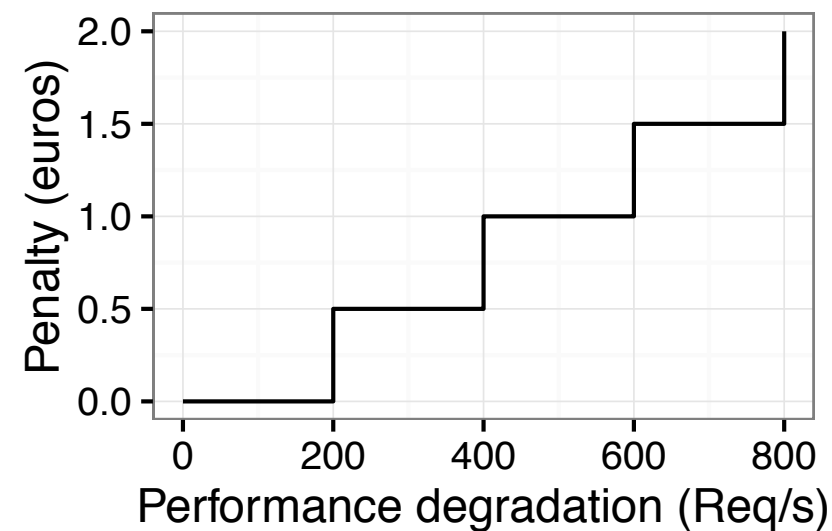
# Hewlett Packard Enterprise



6 to 20 moonshot  
cartridges




SLO




penalty model

| Application | Performance       | Power (W) |
|-------------|-------------------|-----------|
| Website     | 1050 – 3250 Req/s | 360 – 550 |
| G-indexing  | 0 – 565 kPages/h  | 6 – 33    |
| E-indexing  | 0 – 60 kPages/h   | 6 – 33    |


baseline  
(satisfy perf)




carver



“green”  
(max renewable)




# Resulting running costs




# Side tools using simulation

sensitivity analysis  
smart city authority penalties  
SLOs





# seeking for installation sites



# Lesson learned

adaptation requires flexibility at every level  
from hardware to SLOs

workload affinity is a thing

mixing economy and sustainability makes sense (to me)  
good pricing value as a consecutive challenge  
does the energy cost really drive everyone ?



<http://www.dc4cities.eu>