



# Working around Software Defined Networks (SDN) limitations: the compression of flow table & flow management cases

Dino Lopez – [dino.lopez@unice.fr](mailto:dino.lopez@unice.fr)

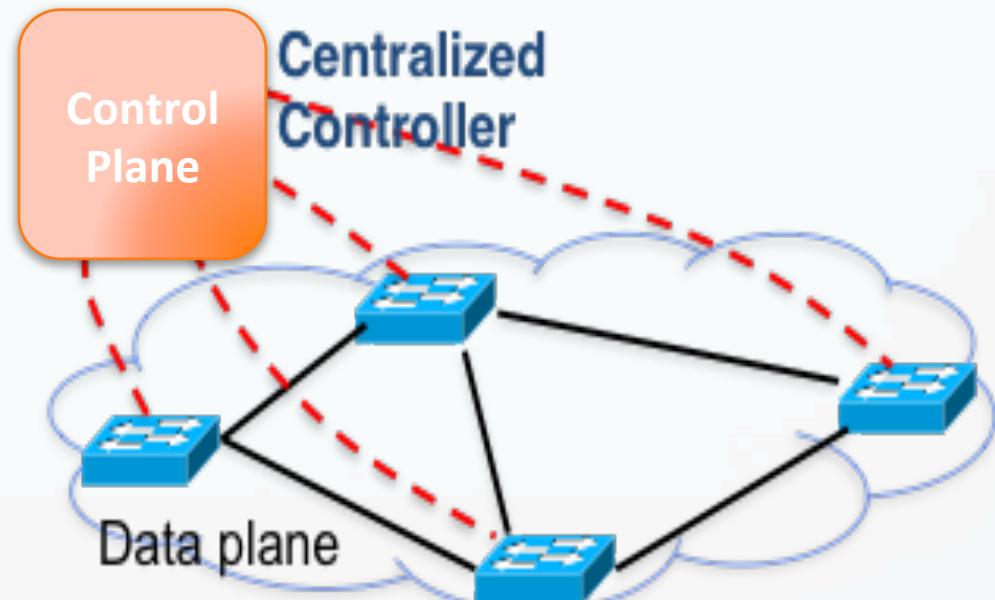
Université Côte d'Azur – Laboratoire I3S/CNRS UMR 7271

SigNet Team

Sophia Antipolis, France

# Motivation

- SDN transforms network administrators into network programmers
  - The control plane and data plane are physically disjoint
  - The controller uses programs to modify the network devices' behavior
- SDN can use flow-based rules to forwards packets
- SDN devices can be hardware-based (e.g. HP switch) or software-based (e.g. OpenVSwitch)



Flow based rules:

Input port	Vlan ID	Vlan PCP	Src. MAC	Dst. MAC	Src. IP	Dst. IP	...	action
------------	---------	----------	----------	----------	---------	---------	-----	--------

12 tuples

# Some challenges

- Richer expressiveness of rules comes at a price
  - need for special hardware → TCAM (**expensive** and **power hungry**)
  - Typical hardware switches support a few thousands rules
  - Ex: our SDN platform, built with an HP 5412zl, supports 3000 hardware rules (up to 65000 software rules)
- The controller-SDN device channel may be the root of new problems
  - Bottleneck
  - Security risks
  - ...



# Improving Network Performance with SDN

- Several propositions based on SDN exist to improve the network performances  
[Popescu2015,Tu2014,Heller10,Cao2014,Han2014,Jain2013,...]
- In our case, we employed SDN to improve the performance in Data Center networks by providing
  - Fine-grain routing with per flow-based rules
  - Short flows protection

# Enabling Fine-grain Routing with SDN

Work in Collaboration between the COATI team (INRIA) and SigNet (I3S)

# How many forwarding rules?

- Up to one billion for the task management in the cloud

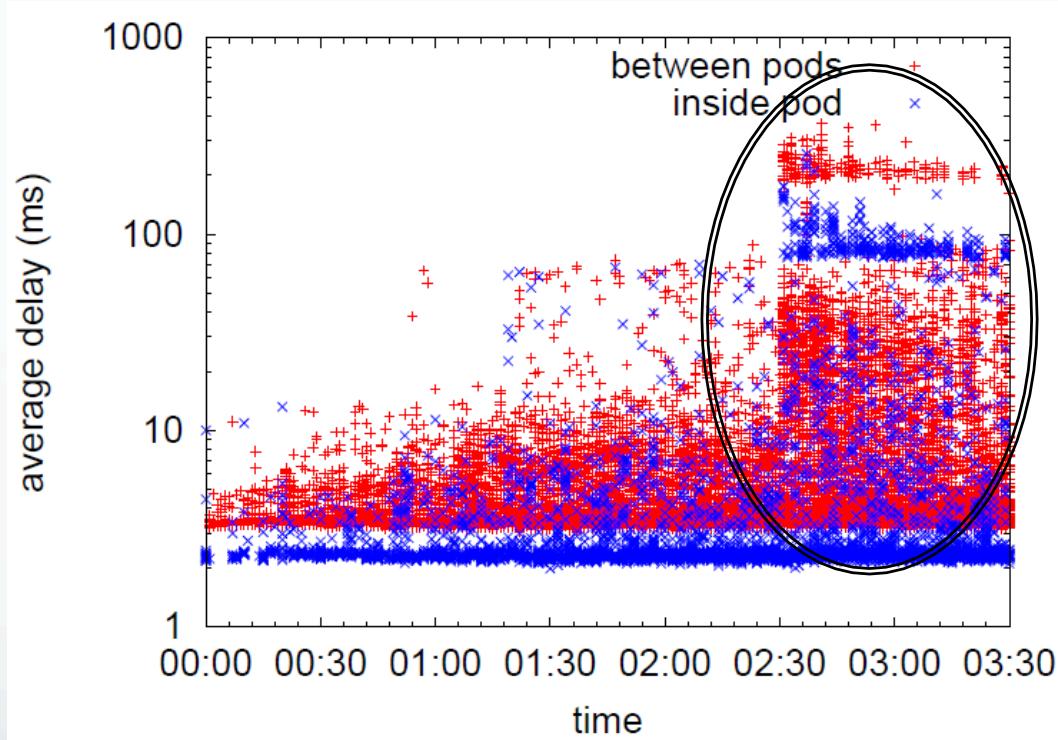
Masoud Moshref, Minlan Yu, Abhishek Sharma, and Ramesh Govin- dan, *vCRIB: Virtualized Rule Management in the Cloud*, USENIX HotCloud (Berkeley, CA, USA), 2012

- Top of the rack up to 78000 rules

Andrew R. Curtis, Jerey C. Mogul, Jean Tourrilhes, Praveen Yala- gandula, Puneet Sharma, and Sujata Banerjee, *DevoFlow: scaling flow management for high-performance networks*, SIGCOMM Comput. Commun. Rev. **41** (2011), no. 4, 254–265

- So what is the impact of limited TCAM ?

# Impact of full TCAM in SDN-based networks



When the routing table is full, the switch needs to contact the controller for every incoming packet from an unknown flow (high delay)

# Number of rules but also rate at which to install them....

- Problem: rule installation processed **by slow path** of SDN routers
  - Can handle a **few hundreds** events per second

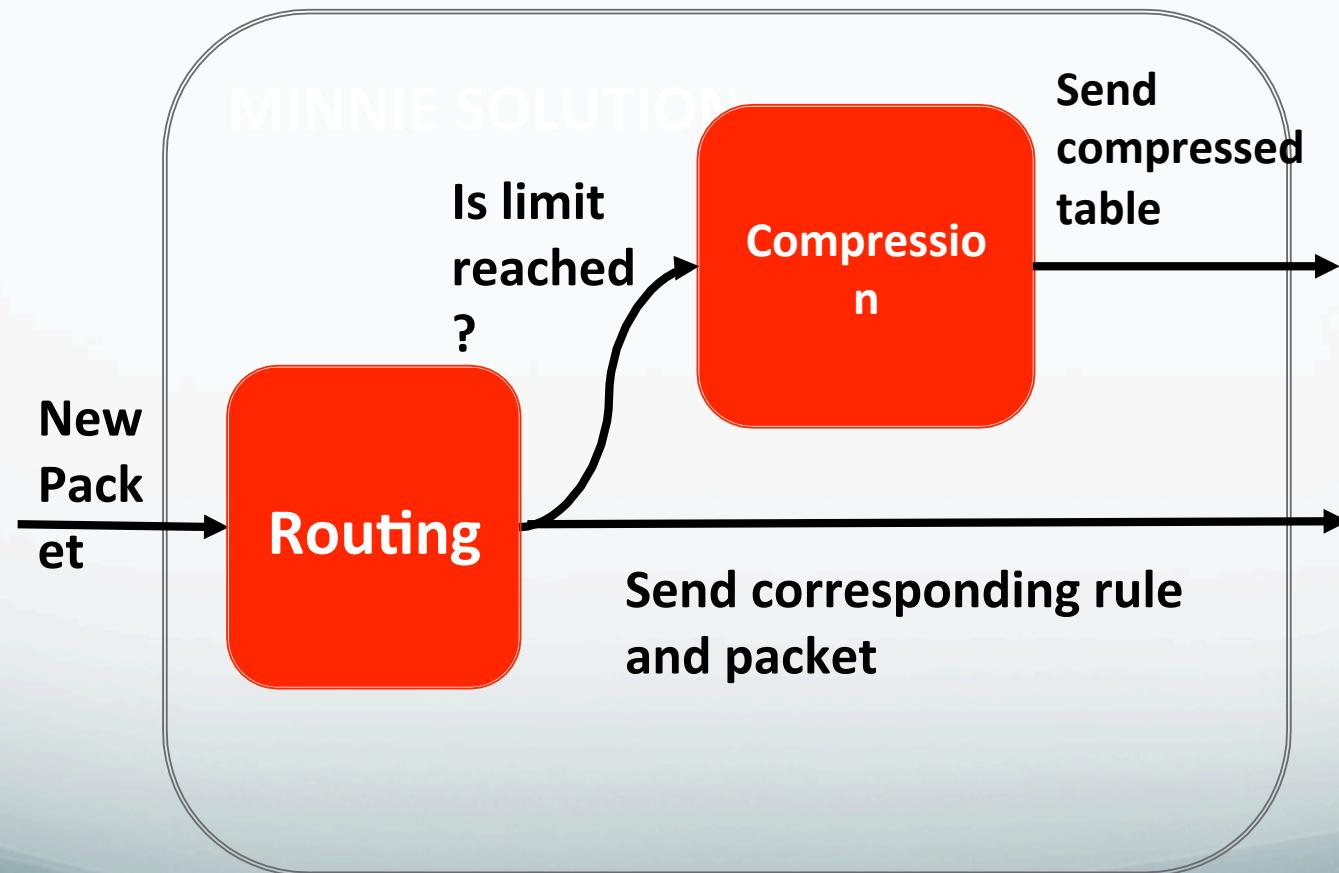
Maciej Kuzniar, Peter Peresíni, Dejan Kostic:  
What You Need to Know About SDN Flow Tables. PAM  
2015

# Related Work

- **Eviction** to remove the least interesting rule (e.g. LRU, [Lee2013])
  - Which is the least interesting rule?
  - Frequent contact with controller
- **Split and distribute the rules** in network (e.g. DIFANE [Yu2010])
  - Needs an accurate view of the traffic matrix
- Use **default paths and deviate** from it when the target is close (e.g. OFFICER [DSauvez2015])
- **Compressing using wildcard rules.**
  - *Our proposition resides in this category*

# The *MINNIE* Solution

Myriana Rifai, Nicolas Huin, Christelle Caillouet, Frédéric Giroire, Dino Lopez Pacheco, Joanna Moulierac, Guillaume Urvoy-Keller. "Too Many SDN Rules? Compress Them with MINNIE". GLOBECOM 2015.



# Minnie: An any-field Compression Solution

Flow	Output port
(0, 4)	Port-4
(0, 5)	Port-5
(0, 6)	Port-5
(1, 4)	Port-6
(1, 5)	Port-4
(1, 6)	Port-6
(2, 4)	Port-4
(2, 5)	Port-5
(2, 6)	Port-6

Flow	Output port
(1, 4)	Port-6
(1, 5)	Port-4
(0, 6)	Port-5
(*, 4)	Port-4
(*, 5)	Port-5
(*, *)	Port-6

Flow	Output port
(0, 4)	Port-4
(1, 5)	Port-4
(2, 4)	Port-4
(2, 5)	Port-5
(0, *)	Port-5
(*, *)	Port-6

Flow	Output port
(0, 5)	Port-5
(0, 6)	Port-5
(1, 4)	Port-6
(1, 6)	Port-6
(2, 5)	Port-5
(2, 6)	Port-6
(*, *)	Port-4

(a) Original routing table

(b) Comp by source

(c) Comp by destination

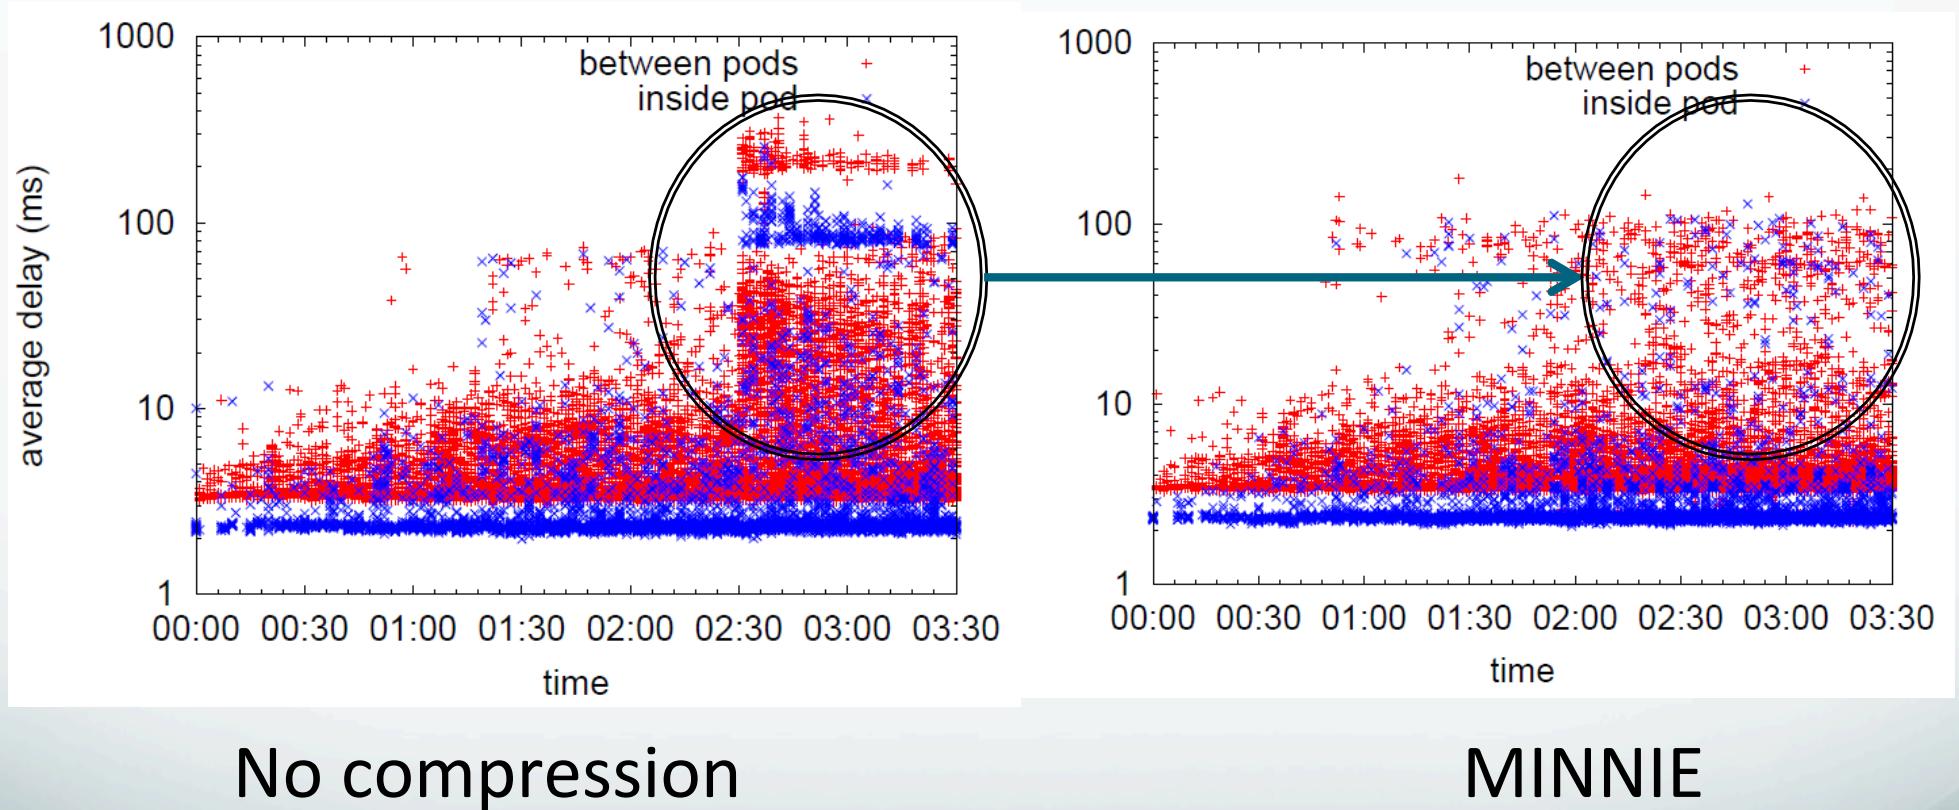
(d) Best default route

The smallest forwarding table is chosen

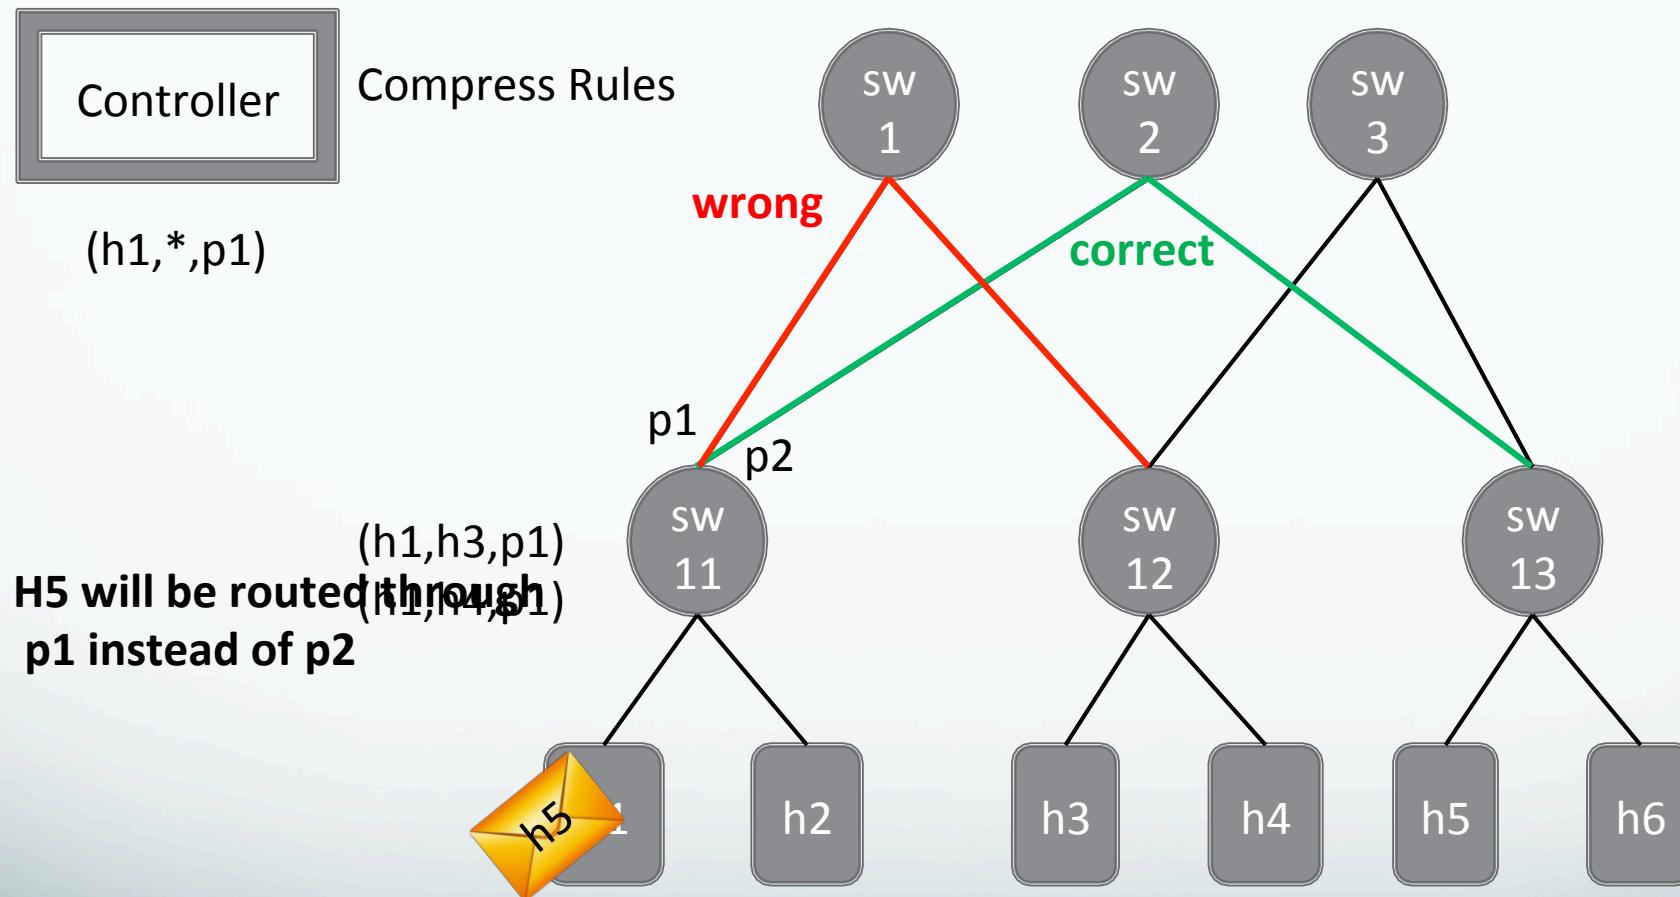
# Evaluating MINNIE

- MINNIE has been evaluated by both simulations and experiments in a testbed
- Same code used during the simulation phase was used for the experimental part
  - We deployed MINNIE on top of a Beacon-based controller
- We could evaluate several parameters
  - Compression rate for large and most famous DC topologies : simulation
  - Impact of compression over the network traffic (delay & jitter): experiments on a testbed
- More information available at [Rifai2015]

# MINNIE in action



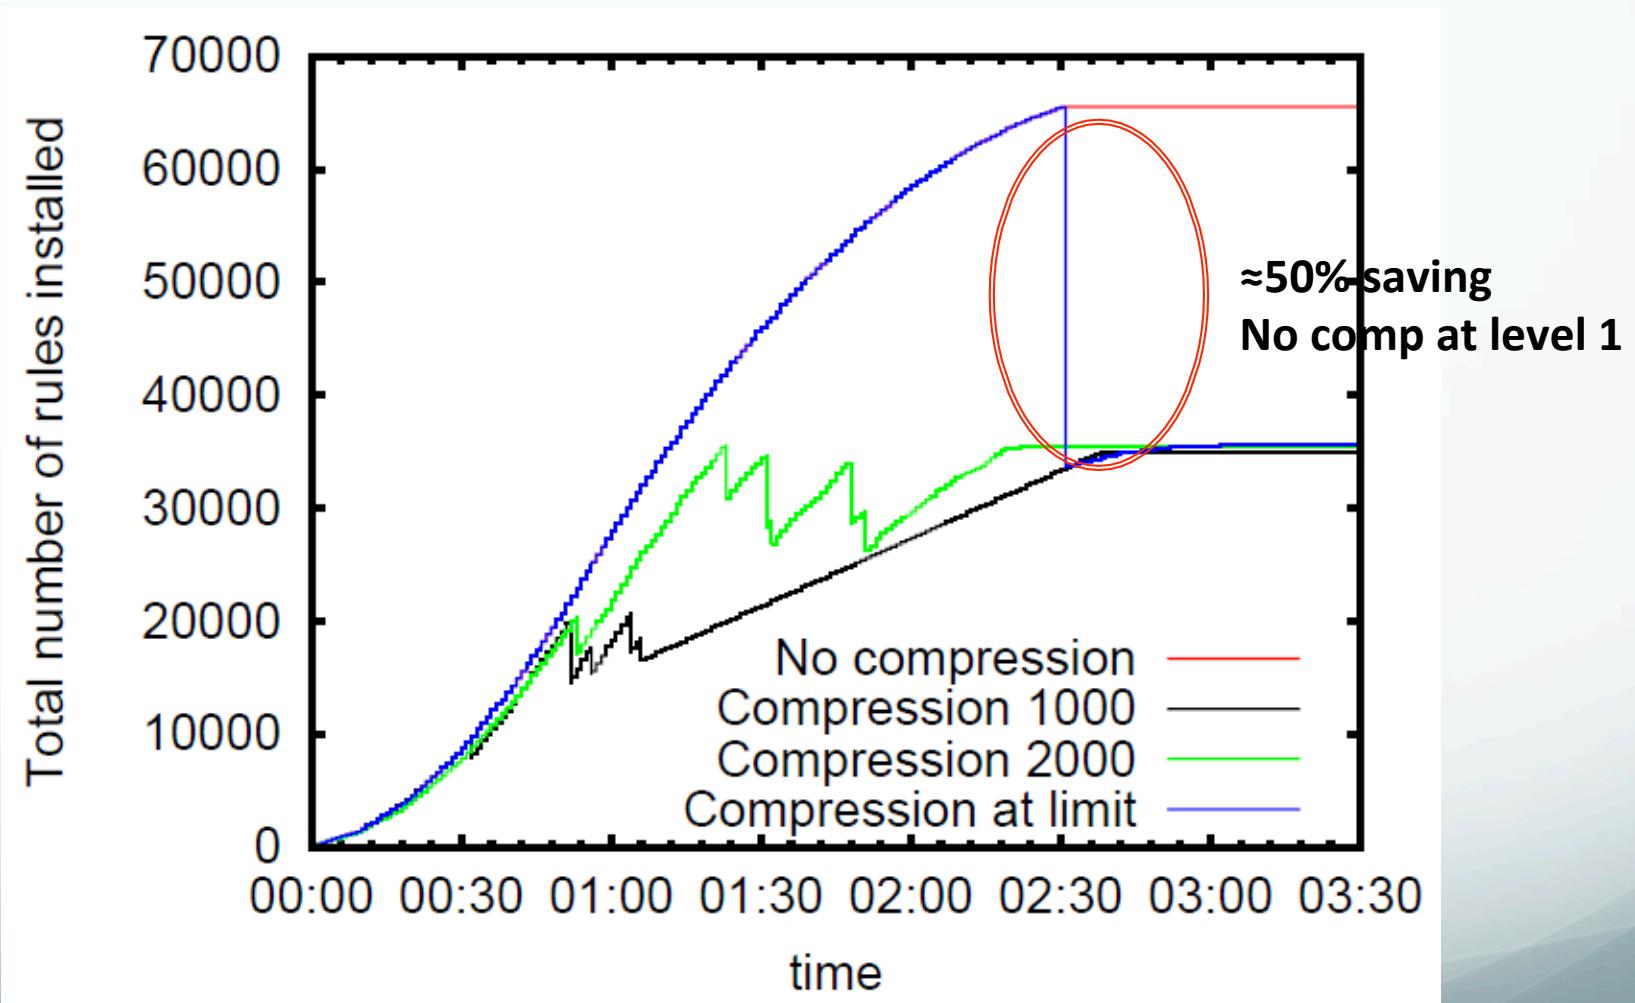
# But... can we safely compress the routing table at the access switches?



# Results

## k-4 pods Fat Tree topology

Total number of rules installed on the physical switch



# Short-Lived flows protection

# Motivation

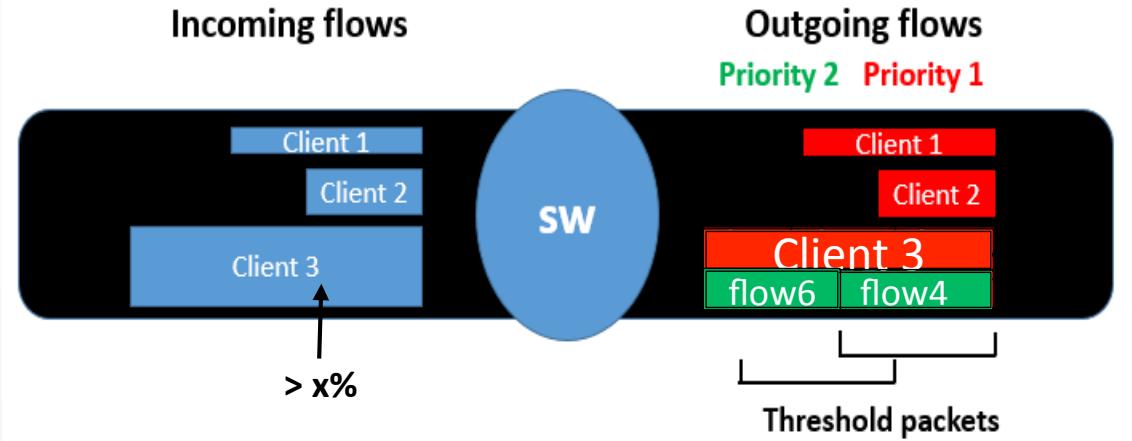
- Large flows represent at least a least percentage of network connections [Benson2010]
- These large flows tend to penalize short flows by consuming all network resources:
  - Majority of network connections are penalized ( latency, losses, low QoE)

# Related Work

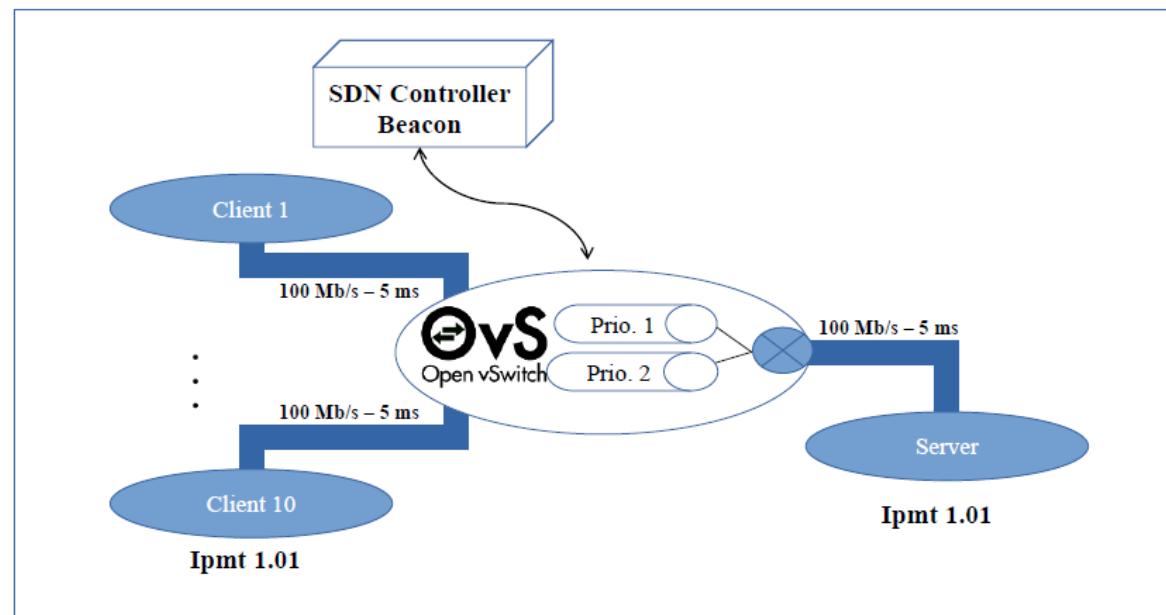
- Use longer rarely used paths for large flows to decrease completion time [Tso2013]
- Detect long flows at the servers and find the most adapted path for such flows [Curtis2011]
- Active queue management strategies: RED, CoDel, etc.
- TCP modifications + multi-bit network signaling: DCTCP

# Coarse-grained Scheduling with Software-Defined Networking Switches

- The scheduler monitors and *retrieves* per rules statistics every  $T$  ms
- State-full scheduler:
    - Install one rule per flow
    - Use the lower priority queue after threshold\_pkts sent by the flow
  - Scalable scheduler:
    - One rule → several flows
    - When such a rule uses more than  $x\%$  of link capacity, install rules per flow
    - Monitor the new per flow rules
    - Use the lower priority queue after threshold\_pkts sent by the flow
    - After few Tmonitor cycles, large flows are isolated and the grouped forwarding rule is reinstalled



# Experimental testbed with Mininet

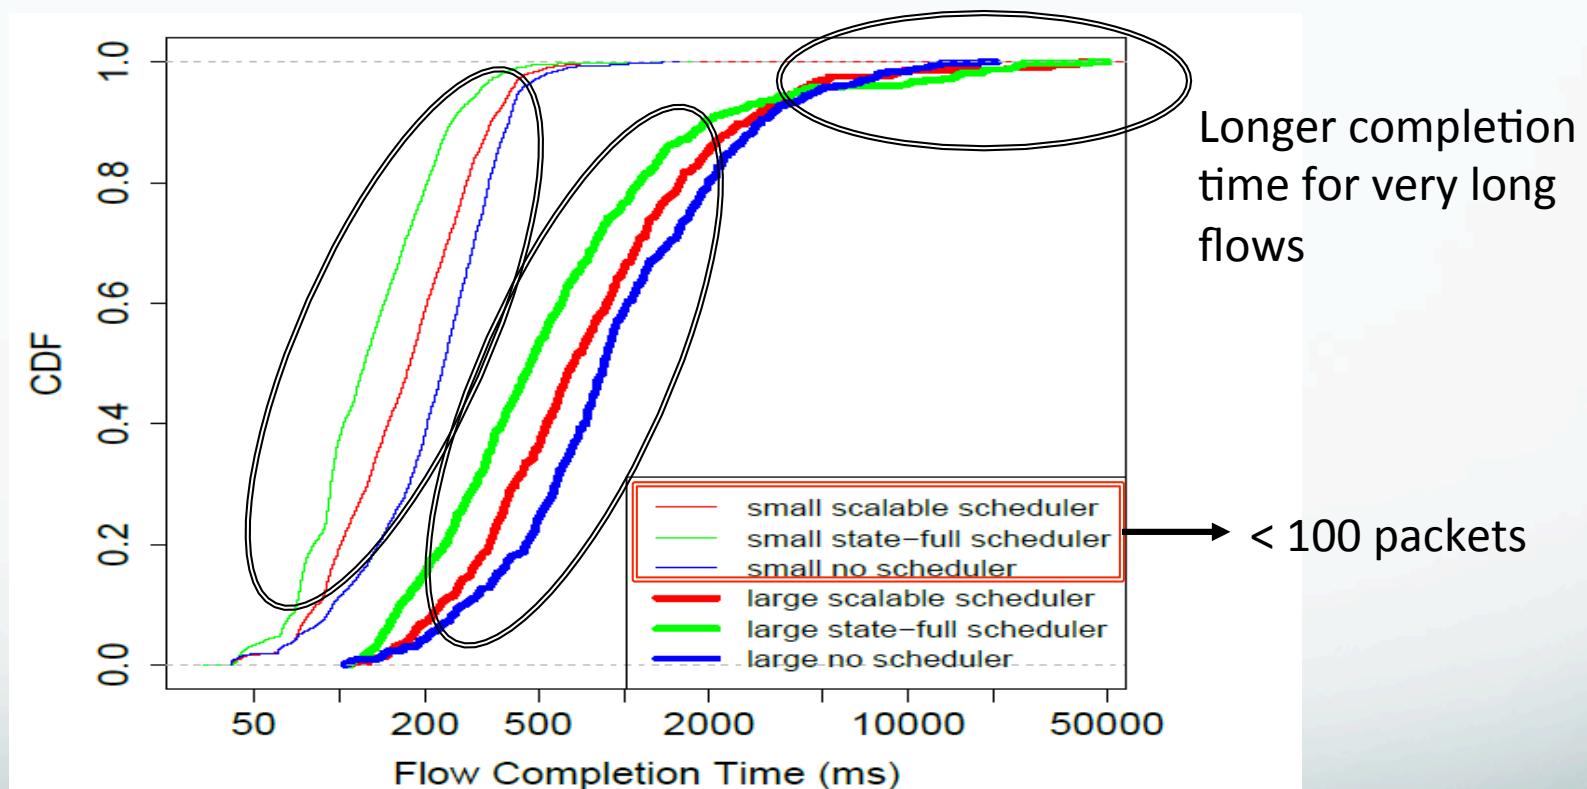


Network traffic characteristics:

- Zipf distribution with a flow size between 15KB (10 packets) and 10 MB
- The average flow size is around 100 packets
- Short flows represents, in average, 90% of the total number of flows

# Results

Flow completion time is better in statefull and scalable scheduler for short and medium long flows



Flow completion time CDF for long and short flows

# Return on experience & conclusions

# Technical Challenges

- Are SDN devices mature enough? Some actions required by the OpenFlow standard are not supported in ASICs by some SDN hardware models. Ex.
  - MAC address changes
  - Port matching (transport header dissection?)
  - Per-flow statistics retrieval
- SDN devices are completely reactive. Is it the best design choice?
  - Every new action (forward, drop, ...) must be explicitly indicated. Can it be improved? Push up some “intelligence”?
  - Statistics retrieval needs request/response messages. Should it be part of the OF-Mon (Notification Framework)?
- Minor technical problems
  - Some switches model does not support queuing priority

# Conclusions

- Our experiences show that SDN can improve the network performance
  - Current SDN hardware devices can indeed provide fine grain network management
  - Monitoring for management can be done by leveraging SDN features
- From the research perspective
  - SDN catalyzes the research reproducibility
  - SDN solutions can be more easily tested and improved
    - SDN hardware can be used to test your proposition in real environments
    - SDN software switches are really useful in absence of hardware devices
- It is now in your hands: deploy your solution
  - Make SDN evolve in the right way